

GROUP BY and HAVING: Solutions

Schema

Student(sID, surName, firstName, campus, email, cgpa) Offering[dept, cNum] \subseteq Course[dept, cNum]
Course(dept, cNum, name, breadth) Took[sID] \subseteq Student[sID]
Offering(oID, dept, cNum, term, instructor) Took[oID] \subseteq Offering[oID]
Took(sID, oID, grade)

Questions

1. Write a query to find the average grade, minimum grade, and maximum grade for each offering.

Solution:

```
select avg(grade), min(grade), max(grade)
from Took
group by oid;
```

Output:

avg	min	max
59.0000000000000000	39	98
60.6666666666666667	45	75
70.5000000000000000	52	89
. . . rows omitted		
75.0000000000000000	54	96
78.0000000000000000	78	78
83.0000000000000000	71	89

(23 rows)
(1 row)

2. Which of these queries is legal?

```
SELECT surname, sid
FROM Student, Took
WHERE Student.sid = Took.sid
GROUP BY sid;
```

```
SELECT surname, Student.sid
FROM Student, Took
WHERE Student.sid = Took.sid
GROUP BY campus;
```

```
SELECT instructor, max(grade),
count(Took.oid)
FROM Took, Offering
WHERE Took.oid = Offering.oid
GROUP BY instructor;
```

```
SELECT Course.dept, Course.cnum,
count(oid), count(instructor)
FROM Course, Offering
WHERE Course.dept = Offering.dept and
Course.cnum = Offering.cnum
GROUP BY Course.dept, Course.cnum
ORDER BY count(oid);
```

Solution: Here's the result of each:

```
SELECT surname, sid
FROM Student, Took
WHERE Student.sid = Took.sid
GROUP BY sid;
ERROR: column reference "sid" is ambiguous
LINE 1: SELECT surname, sid
                  ^
```

instructor	max	count
Heap	82	1
Miller	91	1
Johancsik	99	3
. . . etc.		
Mylopoulos	96	3
Percy	98	4
Mendel	75	3
(17 rows)		

```
ERROR: column "student.surname" must
appear in the GROUP BY clause or be used in
an aggregate function
LINE 1: SELECT surname, Student.sid
                  ^
```

dept	cnum	count	count
ENV	200	1	1
. . . etc.			
CSC	263	3	3
CSC	148	4	4
CSC	207	4	4
CSC	343	5	5
(17 rows)			

3. Find the sid and minimum grade of each student with an average over 80.

Solution:

```
SELECT SID, min(grade)
FROM Took
GROUP BY SID
HAVING AVG(grade) > 80;
```

Output:

sid	min
98000	54
99999	52
(2 rows)	

4. Find the sid, surname, and average grade of each student, but keep the data only for those students who have taken at least 10 courses.

Solution:

```
SELECT Student.sID, surname, avg(grade)
FROM Student, Took
WHERE Student.sID = Took.sID
GROUP BY Student.sID
HAVING count(grade) >= 10;
```

Output:

sid	surname	avg
98000	Fairgrieve	83.2000000000000000
99999	Ali	84.5833333333333333
157	Lakemeyer	75.9333333333333333

(3 rows)

5. For each student who has passed at least 10 courses, report their sid and average grade on the courses that they passed.

Solution:

```
SELECT sid, AVG(grade)
FROM took
WHERE grade >= 50
GROUP BY sid
HAVING count(*) >= 10;
```

Output:

sid	avg
98000	83.2000000000000000
99999	84.5833333333333333
157	78.5714285714285714

(3 rows)

There is a lot going on here. Be sure you are clear on the difference between WHERE and HAVING, and which rows are left at the moment where the HAVING condition is checked for each group.

6. For each student who has passed at least 10 courses, report their sid and average grade on *all of their courses*.

Solution: Here, because we want a filter applied (only passing grades count) when choosing which students to report on, but we don't want that filter applied when we compute their average grade. A single query, with a single WHERE clause, can't accomplish this. Views to the rescue!

```
CREATE VIEW Seniors AS
SELECT sid
FROM Took
WHERE grade >= 50
GROUP BY sid
HAVING count(*) >= 10;
```

```
SELECT Seniors.sid, AVG(grade)
FROM Seniors, Took
WHERE seniors.sid = Took.sid
GROUP BY Seniors.sid;
```

Output:

```
sid | avg
-----+-----
98000 | 83.2000000000000000
99999 | 84.5833333333333333
157 | 75.9333333333333333
(3 rows)
```

Notice that the average for student 157 is different than it was in the previous question. This is because that student failed one course, and it now is allowed to pull down the reported average.

7. Which of these queries is legal?

```
SELECT dept
FROM Took, Offering
WHERE Took.oID = Offering.oID
GROUP BY dept
HAVING avg(grade) > 75;
```

```
SELECT Took.oID, dept, cNum, avg(grade)
FROM Took, Offering
WHERE Took.oID = Offering.oID
GROUP BY Took.oID
HAVING avg(grade) > 75;
```

```
SELECT Took.oID, avg(grade)
FROM Took, Offering
WHERE Took.oID = Offering.oID
GROUP BY Took.oID
HAVING avg(grade) > 75;
```

```
SELECT oID, avg(grade)
FROM Took
GROUP BY sID
HAVING avg(grade) > 75;
```

Solution: Here's the result of each:

```
dept
-----
EEB
ANT
HIS
CSC
(4 rows)
```

```
ERROR: column "offering.dept" must appear
in the GROUP BY clause or be used in an
aggregate function
LINE 1: SELECT Took.oID, dept,
        cNum, avg(grade)
```

```
oid | avg
-----+-----
8 | 92.0000000000000000
28 | 91.0000000000000000
. . . rows omitted
7 | 83.0000000000000000
(11 rows)
```

```
ERROR: column "took.oid" must appear
in the GROUP BY clause or be used in an
aggregate function
LINE 1: SELECT oID, avg(grade)
```